

Shariq Mukadam

+91-7378850558 | shariqmukadam.work@gmail.com | [linkedin.com/in/shariq-mukadam](https://www.linkedin.com/in/shariq-mukadam) | github.com/lamShariqMukadam

Education

Bharati Vidyapeeth Deemed to be University (BVDU)

Bachelor of Computer Applications (BCA) – CGPA: 8.61/10

Aug 2023 – May 2026

Pune, India

Technical Skills

Languages: Python, Java, SQL, JavaScript

AI & ML: PyTorch, Scikit-learn, TensorFlow, CLIP, Sentence-Transformers, LightGBM, SHAP

LLMs & Agents: LangChain, LangGraph, Claude API, Groq API, Ollama, Whisper

Data & Scraping: Pandas, NumPy, Selenium, BeautifulSoup, OpenPyXL, pdfplumber

Web & APIs: FastAPI, React.js, Node.js, HTML, CSS, Streamlit

Databases: MySQL, Pinecone, PostgreSQL

Tools & Platforms: Git, GitHub, Docker, VS Code, Jupyter Notebook, Google Colab, Hugging Face, Vercel

Projects

FreelanceContractAgent – Freelance Contract Negotiation Agent | LangGraph, pdfplumber, FastAPI, Streamlit

- Built a 4-agent Agentic AI system via LangGraph that classifies contract clauses, generates counter-proposals for unfavorable terms, monitors scope creep, and compiles dispute evidence files autonomously.
- 56M Indian freelancers lack legal support and 40% face payment disputes; the Agentic AI pipeline replaces 2 lakh lawyer fees by handling the full contract lifecycle from signing to dispute resolution.
- Covers all three critical phases – pre-signing negotiation, mid-project scope protection, and post-dispute evidence organization – in a single Agentic AI pipeline via a simple file upload interface.

ReturnSight – E-Commerce Return Probability Predictor | LightGBM, PyTorch, SHAP, FastAPI, React

[GitHub](#)

- Built a multi-modal AI pipeline fusing product review text via sentence-transformers, and seller history via LightGBM – generating a per-listing return probability score via FastAPI with 200–500ms CPU inference.
- E-commerce returns cost \$890B globally in 2024; the pipeline assigns a 0–100% risk score per listing with SHAP explainability pinpointing whether review mismatches, rating anomalies, or price signals drive the risk.
- Deployed as a real-time scoring API providing instant return-risk assessment with SHAP feature attribution; achieved AUC-ROC 0.8302 and Average Precision 0.7022 on 91K held-out test products with zero risk-category flips across 15 real-world validation listings – enabling sellers to fix high-risk listings before generating losses.

JobHarvester – Job Market Intelligence Agent | Selenium, BeautifulSoup, LangChain, OpenPyXL

[GitHub](#)

- Built an AI Agent scraping 200+ live job postings across LinkedIn, Naukri, and Internshala via Selenium and BeautifulSoup, passing each through an LLM extraction layer to output structured skill frequency data in a multi-sheet Excel report.
- AI role JDs show 70%+ inconsistency across companies causing candidates to waste 6–12 months on wrong skills; the agent clusters postings by company stage producing a heatmap of what actually matters versus JD filler.
- Personalization layer compares a user's skill profile against the most relevant job cluster and outputs a prioritized gap analysis – a research report that would otherwise take weeks to compile manually.

Certifications

Python Data Associate – DataCamp

Jun 2026

SQL Associate – DataCamp

Jun 2026

Data Structures and Algorithms using Java – Infosys Springboard

Sep 2025

Problem Solving (Intermediate) – HackerRank

Sep 2025